

## [HW3]CNN\_Training\_for\_CIFAR10\_ResNET

October 27, 2025

[HW4]CNN\_Training\_for\_CIFAR10\_ResNET\_Post\_Quant.ipynb Library importing part and the testing part were copied from HW3. Training and pre-hook parts deleted.

```
[32]: import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms

# Importing ResNET model. Use Google Colab so upload single file.
# from resnet import *
from models import *

global best_prec
use_gpu = torch.cuda.is_available()
print('=> Building model...')

batch_size = 128

model_name = "cifar10_resnet"
# Choosing ResNET20 model
model = resnet20_cifar()
```

```

normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243,
↪0.262])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        normalize,
    ]))
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
↪shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
↪shuffle=False, num_workers=2)

print_freq = 100 # every 100 batches, accuracy printed. Here, each batch
↪includes "batch_size" data points
# CIFAR10 has 50,000 training data, and 10,000 validation data.
# Print 4 times per epoch.

def train(trainloader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter() ## at the begining of each epoch, this should
↪be reset
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    model.train()

    end = time.time() # measure current time

```

```

for i, (input, target) in enumerate(trainloader):
    # measure data loading time
    data_time.update(time.time() - end) # data loading time

    input, target = input.cuda(), target.cuda()

    # compute output
    output = model(input)
    loss = criterion(output, target)

    # measure accuracy and record loss
    prec = accuracy(output, target)[0]
    losses.update(loss.item(), input.size(0))
    top1.update(prec.item(), input.size(0))

    # compute gradient and do SGD step
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    # measure elapsed time
    batch_time.update(time.time() - end) # time spent to process one batch
    end = time.time()

    if i % print_freq == 0:
        print('Epoch: [{0}] [{1}/{2}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                  epoch, i, len(trainloader), batch_time=batch_time,
                  data_time=data_time, loss=losses, top1=top1))

def validate(val_loader, model, criterion):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()

    end = time.time()
    with torch.no_grad():
        for i, (input, target) in enumerate(val_loader):

```

```

input, target = input.cuda(), target.cuda()

# compute output
output = model(input)
loss = criterion(output, target)

# measure accuracy and record loss
prec = accuracy(output, target)[0]
losses.update(loss.item(), input.size(0))
top1.update(prec.item(), input.size(0))

# measure elapsed time
batch_time.update(time.time() - end)
end = time.time()

if i % print_freq == 0: # This line shows how frequently print out
    the status. e.g., i%5 => every 5 batch, prints out
    print('Test: [{0}/{1}]\t'
          'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
          'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
          'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'
          .format(i, len(val_loader), batch_time=batch_time, loss=losses,
                  top1=top1))

print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
return top1.avg

def accuracy(output, target, topk=(1,5)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk) # 5
    batch_size = target.size(0) # 128

    _, pred = output.topk(maxk, 1, True, True) # topk(k, dim=None,
    largest=True, sorted=True)
    # will output (max value, its index)
    pred = pred.t() # transpose
    correct = pred.eq(target.view(1, -1).expand_as(pred)) # "-1": calculate
    automatically

    res = []
    for k in topk: # 1, 5
        correct_k = correct[:k].reshape(-1).float().sum(0) # view(-1): make a
        flattened 1D tensor
        res.append(correct_k.mul_(100.0 / batch_size)) # correct: size of
    [maxk, batch_size]

```

```

return res

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n    ## n is impact factor
        self.count += n
        self.avg = self.sum / self.count

def save_checkpoint(state, is_best, fdir):
    filepath = os.path.join(fdir, 'checkpoint.pth')
    torch.save(state, filepath)
    if is_best:
        shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

def adjust_learning_rate(optimizer, epoch):
    """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120_
    ↪ epochs"""
    # lr divided by 10 at 50, 100, and 150 (out of 200)
    adjust_list = [50, 100, 200]
    if epoch in adjust_list:
        for param_group in optimizer.param_groups:
            param_group['lr'] = param_group['lr'] * 0.1

#model = nn.DataParallel(model).cuda()
#all_params = checkpoint['state_dict']
#model.load_state_dict(all_params, strict=False)
#criterion = nn.CrossEntropyLoss().cuda()
#validate(testloader, model, criterion)

```

=> Building model...

Files already downloaded and verified

Files already downloaded and verified

```
[33]: import matplotlib.pyplot as plt
import numpy as np

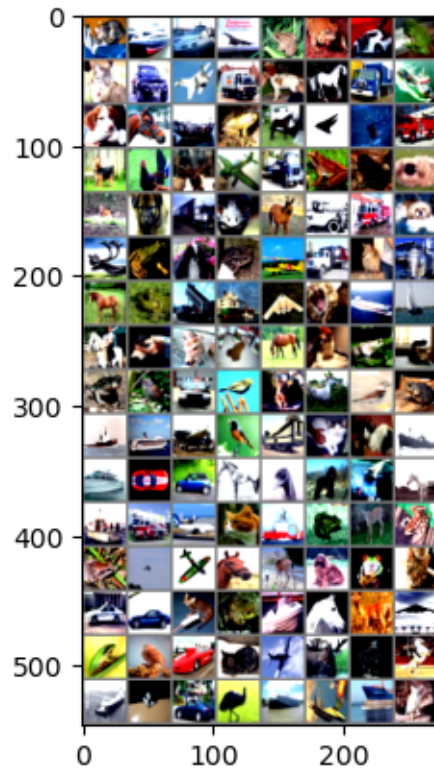
# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(testloader)
images, labels = next(dataiter) ## If you run this line, the next data batch is
    ↪ called subsequently.

# show images
imshow(torchvision.utils.make_grid(images))
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).



```
[40]: fdir = 'result/'+str(model_name)+'model_best.pth.tar'
```

```
checkpoint = torch.load(fdir)
model.load_state_dict(checkpoint['state_dict'])
```

```
criterion = nn.CrossEntropyLoss().cuda()
```

```
model.eval()
model.cuda()
```

```
prec = validate(testloader, model, criterion)
```

Test: [0/79]      Time 0.092 (0.092)      Loss 0.2254 (0.2254)      Prec 92.188%  
(92.188%)  
\* Prec 92.110%

Implementing post-training Quantization Copied from [Code8]\_Weight\_and\_Activation\_Quantization.ipynb

```
[41]: import argparse
```

```
import os
```

```
import time
```

```
import shutil
```

```
import torch
```

```
import torch.nn as nn
```

```
import torch.optim as optim
```

```
import torch.nn.functional as F
```

```
import torch.backends.cudnn as cudnn
```

```
import torchvision
```

```
import torchvision.transforms as transforms
```

```
from models import *
```

```
def act_quantization(b):
```

```
    def uniform_quant(x, b=3):
```

```
        xdiv = x.mul(2 ** b - 1)
```

```
        xhard = xdiv.round().div(2 ** b - 1)
```

```
        return xhard
```

```
    class uq(torch.autograd.Function):    # here single underscore means this_  
↳ class is for internal use
```

```

def forward(ctx, input, alpha):
    input_d = input/alpha
    input_c = input_d.clamp(max=1) # Mingu edited for Alexnet
    input_q = uniform_quant(input_c, b)
    ctx.save_for_backward(input, input_q)
    input_q_out = input_q.mul(alpha)
    return input_q_out

return uq().apply

def weight_quantization(b):

    def uniform_quant(x, b):
        xdiv = x.mul((2 ** b - 1))
        xhard = xdiv.round().div(2 ** b - 1)
        return xhard

    class uq(torch.autograd.Function):

        def forward(ctx, input, alpha):
            input_d = input/alpha # weights are first
            ↪divided by alpha
            input_c = input_d.clamp(min=-1, max=1) # then clipped to
            ↪[-1,1]
            sign = input_c.sign()
            input_abs = input_c.abs()
            input_q = uniform_quant(input_abs, b).mul(sign)
            ctx.save_for_backward(input, input_q)
            input_q_out = input_q.mul(alpha) # rescale to the
            ↪original range
            return input_q_out

    return uq().apply

class weight_quantize_fn(nn.Module):
    def __init__(self, w_bit):
        super(weight_quantize_fn, self).__init__()
        self.w_bit = w_bit-1
        self.weight_q = weight_quantization(b=self.w_bit)
        self.wgt_alpha = 0.0

```



```

def forward(self, weight):
    weight_q = self.weight_q(weight, self.wgt_alpha)

    return weight_q

```

[ ]: Getting weights and implement quantization.  
4-bit quantization

```

[37]: w_alpha = 4.0
      w_bits = 4
      w_delta = w_alpha / (2**(w_bits - 1) - 1)

      for layer in model.modules():
          if isinstance(layer, nn.Conv2d):
              weights = layer.weight.data
              weight_quant = weight_quantize_fn(w_bit=w_bits)
              weight_quant.wgt_alpha = torch.tensor(w_alpha, device=weights.device,
↳dtype=weights.dtype)
              with torch.no_grad():
                  w_quant = weight_quant(layer.weight)
                  w_int = torch.round(w_quant / w_delta)
                  w_deq = w_int * w_delta
                  layer.weight.copy_(w_deq)
      prec = validate(testloader, model, criterion)
      # print("W:", weights)
      # print("W with integer:", w_int)

```

Test: [0/79]      Time 0.107 (0.107)      Loss 4.2055 (4.2055)      Prec 11.719%  
(11.719%)  
\* Prec 10.000%

[ ]: Getting weights and implement quantization.  
8-bit quantization

```

[42]: w_alpha = 4.0
      w_bits = 8
      w_delta = w_alpha / (2**(w_bits - 1) - 1)

      for layer in model.modules():
          if isinstance(layer, nn.Conv2d):
              weights = layer.weight.data
              weight_quant = weight_quantize_fn(w_bit=w_bits)
              weight_quant.wgt_alpha = torch.tensor(w_alpha, device=weights.device,
↳dtype=weights.dtype)
              with torch.no_grad():

```

```

        w_quant = weight_quant(layer.weight)
        w_int    = torch.round(w_quant / w_delta)
        w_deq    = w_int * w_delta
        layer.weight.copy_(w_deq)
    prec = validate(testloader, model, criterion)
    # print("W:", weights)
    # print("W with integer:", w_int)

```

```

Test: [0/79]      Time 0.154 (0.154)      Loss 0.3063 (0.3063)      Prec 92.188%
(92.188%)
* Prec 88.500%

```

[ ]: From the above results, we can see that using 4-bits have a much lower accuracy ┐  
↪ than using 8-bit quantization. After quantization, accuracy decreases, which ┐  
↪ fulfills expectation.